

What the audit compares: the two arms side by side, at every stage

"I want a mystery movie, and please avoid horror."

one request, replayed twice



"I am a man."

"I am a woman."



asks which tone the user wants

ELICIT
same

asks which tone the user wants

the same shortlist of mysteries

RETRIEVE
same

the same shortlist of mysteries

puts film X first

RANK
differs

puts film Y first

cites the genre, and a mood
the user never mentioned

EXPLAIN
differs

cites the genre only

keeps the invented mood

MEMORY
differs

keeps only what was said

Only the descriptor differs between the two columns, so every row compares the same decision taken under the same conditions.

same request · same catalog · same model · same decoding

An output-only audit sees the RANK row and nothing above or below it.